


Calibrated model cascade and empirical Pareto frontier across model rungs

(a) Cost–reliability frontier, held-out test split (n=30)



(b) P6 replicate, infra-excluded test split (n=107)

