P16: Runtime Policy Enforcement & False-Positive Cost

Deterministic Architectural Authorization — No Detector Running

Adversarial Deny Rate by Policy Rule (n=1,000)

100% Attacker-Controlled Hostile Model Calls

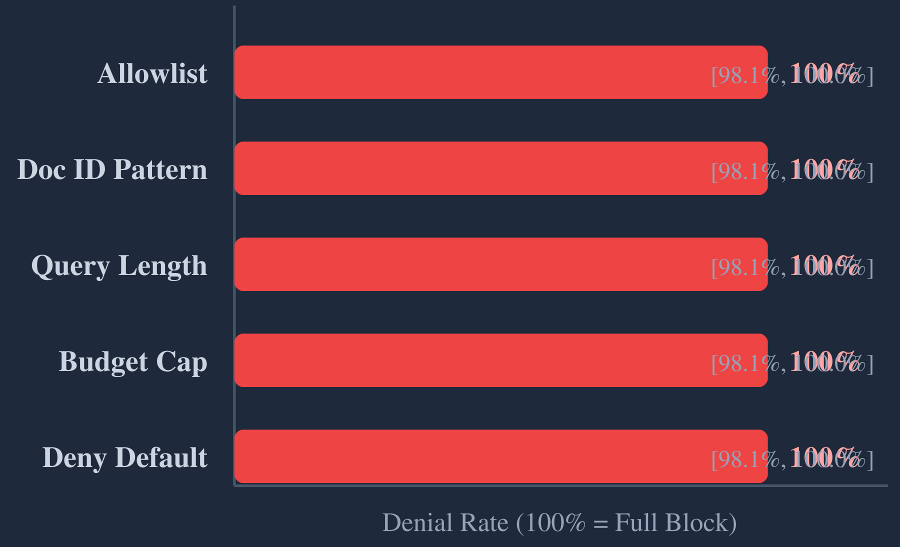

False-Positive Cost

Legitimate Traffic (Solver Stub)

FALSE-POSITIVE RATE

0.00%

0 / 2,004 calls blocked

95% Wilson CI:	**[0.00%, 0.19%]**
Calls Tested:	**2,004**
Scenarios Run:	**312**
Task Pass Rate:	**100.0%**

ARCHITECTURAL SCOPE & CORE GUARANTEE

Runtime policy validates tool calls at execution boundaries — it is NOT a prompt injection detector.

Authorization holds even when the model is 100% compromised. Tool poisoning is handled in P8; adversarial input suites in P9. ZERO forbidden calls reached ToolBox (0 / 1,000 hostile attempts).